

Theoretical Foundations of Buffer Stock Saving

September 6, 2021

Christopher D. Carroll¹

Abstract

This paper builds foundations for rigorous and intuitive understanding of ‘buffer stock’ saving models (?-like models with a wealth target), pairing each theoretical result with quantitative illustrations. After describing conditions under which a consumption function exists, the paper articulates stricter ‘Growth Impatience’ conditions that guarantee alternative forms of stability — either at the population level, or for individual consumers. Together, the numerical tools and analytical results constitute a comprehensive toolkit for understanding buffer stock models.

Keywords Precautionary saving, buffer stock saving, marginal propensity to consume, permanent income hypothesis, income fluctuation problem

JEL codes D81, D91, E21

Powered by Econ-ARK

Dashboard: <https://econ-ark.org/materials/BufferStockTheory?dashboard>
REMARK: <https://econ-ark.org/materials/BufferStockTheory>
html: <https://econ-ark.github.io/BufferStockTheory/>
PDF: <https://econ-ark.github.io/BufferStockTheory/BufferStockTheory.pdf>
Slides: <https://econ-ark.github.io/BufferStockTheory/BufferStockTheory-Slides.pdf>
Appendix: <https://econ-ark.github.io/BufferStockTheory#Appendices>
GitHub: <https://github.com/econ-ark/BufferStockTheory>

The dashboard lets users see consequences of alternative parameters in an interactive framework.

The paper’s results can be automatically reproduced using the Econ-ARK/HARK toolkit, which can be cited per our references (?); for reference to the toolkit itself see Acknowledging Econ-ARK. Thanks to the Consumer Financial Protection Bureau for funding the original creation of the Econ-ARK toolkit; and to the Sloan Foundation for funding Econ-ARK’s extensive further development that brought it to the point where it could be used for this project. The toolkit can be cited with its digital object identifier, 10.5281/zenodo.1001067, as is done in the paper’s own references as ?. Thanks to Will Du, James Feigenbaum, Joseph Kaboski, Miles Kimball, Qingyin Ma, Misuzu Otsuka, Damiano Sandri, John Stachurski, Adam Szeidl, Alexis Akira Toda, Metin Uyanik, Mateo Velásquez-Giraldo, Weifeng Wu, Jiaxiong Yao, and Xudong Zheng for comments on earlier versions of this paper, John Boyd for help in applying his weighted contraction mapping theorem, Ryoji Hiraguchi for extraordinary mathematical insight that improved the paper greatly, David Zervos for early guidance to the literature, and participants in a seminar at the Johns Hopkins University, a presentation at the 2009 meetings of the Society of Economic Dynamics for their insights, and at a presentation at the Australian National University.

¹Contact: ccarroll@jhu.edu, Department of Economics, 590 Wyman Hall, Johns Hopkins University, Baltimore, MD 21218, <https://www.econ2.jhu.edu/people/ccarroll>, and National Bureau of Economic Research.

1 Introduction

In the presence of realistic transitory and permanent shocks to income *a la* ? and ?, only one further ingredient is required to construct a microeconomically testable model of consumption: A description of preferences. Zeldes (?) was the first to calibrate a quantitatively plausible example; his paper spawned a literature showing that such models' predictions can match household life cycle data reasonably well, whether or not explicit liquidity constraints are imposed.¹

A connected literature in macroeconomic theory, starting with ?, has derived limiting properties of related infinite-horizon problems, but only in models more complex than the case with just shocks and preferences. The extra complexity has been imposed because standard contraction mapping theorems (beginning with ? and including those building on Stokey et al. (??)) cannot be applied when utility and/or marginal utility are unbounded. Many proof methods also rule out permanent shocks *a la* ?, ?, and ?.²

This paper's first technical contribution is to articulate conditions under which the infinite-horizon version of the Friedman-Muth-Zeldes problem (without complications like a consumption floor or liquidity constraints) defines a contraction mapping whose limiting solution is sensible as the horizon approaches infinity. A 'Finite Value of Autarky Condition' is mostly sufficient (the other imposed condition, the 'Weak Return Impatience Condition',³ is unlikely to bind). Because the infinite horizon solution is the limit of finite-horizon recursions, many intermediate results are also useful for solving finite-horizon problems.

But the paper's main theoretical contribution is to identify, for the infinite-horizon case, conditions under which 'stable' values of the wealth-to-permanent-income ratio exist, either for individual consumers (a consumer's wealth can be predicted to move toward a 'target' ratio) or for the aggregate (the economy as a whole moves toward a 'balanced growth' equilibrium). The requirement for stability is always that the model's parameters satisfy a 'Growth Impatience Condition' whose details depend on the quantity whose stability is of interest. A model that exhibits stability of either kind qualifies as a 'buffer stock' model.⁴

Even without a formal proof of its existence, buffer stock saving has been intuitively understood to underlie central quantitative results in heterogeneous agent macroeconomics; for example, the logic of target saving is central to the claim by ? in the *Handbook of Macroeconomics* that such models explain why, during the Great Recession, middle-class consumers cut their spending more than the poor or the rich. The theory below provides the rigorous basis for this claim: Learning that the future has become

¹See ? or ? for arguments that models with only 'natural' constraints (see below) match a wide variety of facts; for a model with explicit constraints that produces very similar results, see, e.g., ?.

²See the fuller discussion at the end of Section 2.1.

³This is a generalization of a condition in ?.

⁴Such models are neither a subset nor a superset of ? models. But closed economies in which capital results from saving and has declining marginal productivity are always 'buffer stock' economies under some definition of that term, because capital accumulation causes interest rates to fall, which guarantees that a Growth Impatience Condition will hold in equilibrium (see below). The more interesting applications are to populations (or economies) whose marginal saving behavior does not determine the relevant interest rate, or in which the marginal product of capital does not fall as capital is accumulated (again, see below).

more uncertain does not change the urgent imperatives of the poor (their high $u'(c)$ means they — optimally — have little room to maneuver). And, increased labor income uncertainty does not much change the behavior of the rich because it poses little risk to their consumption. Only people in the middle have both the motivation and the wiggle-room to respond by reducing their spending.

Analytical derivations for the proofs also explain many other results familiar from the numerical literature.

The paper begins by defining sufficient conditions for the problem to define a useful (nondegenerate) limiting consumption function (and explains how the model relates to those previously considered). The conditions are interestingly parallel to those required for the **liquidity constrained perfect foresight model**; that parallel is explored and explained. This analysis establishes limiting properties of the consumption function as resources approach infinity, and as they approach their lower bound; using these limits, the contraction mapping theorem is proven.

The next theoretical contribution demonstrates that a corresponding model with an ‘artificial’ liquidity constraint (a model that prohibits borrowing, even by consumers who could certainly repay) is a limiting case of the model without constraints. The analytical appeal of the unconstrained model is that it is both mathematically convenient (e.g., the consumption function is twice continuously differentiable), and arbitrarily close (cf. Section 2.10) to less tractable models. The idea is to prove a proposition in a congenial environment, then to define the analogous proposition as holding (in the limit) if it holds for in the limit as the horizon extends to infinity.

In proving the remaining theorems, the **next section** examines the key properties of the model. First, as **cash approaches infinity** the expected growth rate of consumption and the marginal propensity to consume (MPC) converge to their values in the perfect foresight case. Second, as **cash approaches zero** the expected growth rate of consumption approaches infinity, and the MPC approaches a simple analytical limit. Next, the central theorems articulate conditions under which different measures of ‘growth impatience’ imply useful conclusions about points of stability (‘target’ or ‘balanced growth’ points).

The final section elaborates the conditions under which, even with a fixed aggregate interest rate that differs from the time preference rate, a small open economy populated by buffer stock consumers has a balanced growth equilibrium in which growth rates of consumption, income, and wealth match the exogenous growth rate of permanent income (equivalent, here, to productivity growth). In the terms of ?, buffer stock saving is an appealing method of ‘closing’ a small open economy model, because it requires no ad-hoc assumptions. Not even liquidity constraints.⁵

⁵The paper’s insights are instantiated in the **Econ-ARK** toolkit, whose **buffer stock saving module** flags parametric choices under which a problem is degenerate or under which stable ratios of wealth to income may not exist.

2 The Problem

2.1 Setup

The infinite horizon solution is the (limiting) first-period solution to a sequence of finite-horizon problems as the horizon (the last period of life) becomes arbitrarily distant.

That is, for the value function, fixing a terminal date T , we are interested in the term \mathbf{v}_{T-n} in the sequence of value functions $\{\mathbf{v}_T, \mathbf{v}_{T-1}, \dots, \mathbf{v}_{T-n}\}$. We will say that the problem has a ‘nondegenerate’ infinite horizon solution if, corresponding to that \mathbf{v} , as $n \uparrow \infty$ there is a limiting consumption function $c(m) = \lim_{n \uparrow \infty} c_{T-n}$ which is neither $c(m) = 0$ everywhere (for all m) nor $c(m) = \infty$ everywhere.

Concretely, a consumer born n periods before date T solves the problem

$$\mathbf{v}_{T-n} = \max \mathbb{E}_t \left[\sum_{i=0}^n \beta^i u(\mathbf{c}_{t+i}) \right]$$

where the Constant Relative Risk Aversion (CRRA) utility function

$$u(\bullet) = \bullet^{1-\rho} / (1-\rho) \quad (1)$$

exhibits relative risk aversion $\rho > 1$.⁶ The consumer’s initial condition is defined by market resources \mathbf{m}_t and permanent noncapital income \mathbf{p}_t , which both are positive,

$$\{\mathbf{p}_t, \mathbf{m}_t\} \in (0, \infty), \quad (2)$$

and the consumer cannot die in debt,

$$\mathbf{c}_T \leq \mathbf{m}_T. \quad (3)$$

In the usual treatment, a dynamic budget constraint (DBC) incorporates several elements that jointly determine next period’s \mathbf{m} (given this period’s choices); for the detailed analysis here, it will be useful to disarticulate and describe the steps:

$$\begin{aligned} \mathbf{a}_t &= \mathbf{m}_t - \mathbf{c}_t \\ \mathbf{k}_{t+1} &= \mathbf{a}_t \\ \mathbf{b}_{t+1} &= \mathbf{k}_t R \\ \mathbf{p}_{t+1} &= \mathbf{p}_t \underbrace{\Gamma^{\psi_{t+1}}}_{\equiv \Gamma_{t+1}} \\ \mathbf{m}_{t+1} &= \mathbf{b}_{t+1} + \mathbf{p}_{t+1} \xi_{t+1}, \end{aligned} \quad (4)$$

where \mathbf{a}_t indicates the consumer’s assets at the end of period t , which translate one-for-one into capital \mathbf{k}_{t+1} at the beginning of the next period, which (before the consumption choice) grows by a fixed interest factor $R = (1+r)$, so that \mathbf{b}_{t+1} is the consumer’s financial (‘bank’) balances before next period’s consumption choice;⁷ \mathbf{m}_{t+1} (‘market resources’) is

⁶The main results also hold for logarithmic utility which is the limit as $\rho \rightarrow 1$ but incorporating the logarithmic special case in the proofs is omitted because it would be cumbersome.

⁷Allowing a stochastic interest factor is straightforward but adds little insight for our purposes; however, see ?, ?, and ? for the implications of capital income risk for the distribution of wealth and other interesting questions not considered here.

the sum of financial wealth \mathbf{b}_{t+1} and noncapital income $\mathbf{p}_{t+1}\xi_{t+1}$ (permanent noncapital income \mathbf{p}_{t+1} multiplied by a mean-one iid transitory income shock factor ξ_{t+1} ; transitory shocks are assumed to satisfy $\mathbb{E}_t[\xi_{t+n}] = 1 \forall n \geq 1$). Permanent noncapital income in $t+1$ is equal to its previous value, multiplied by a growth factor Γ , modified by a mean-one iid shock ψ_{t+1} , $\mathbb{E}_t[\psi_{t+n}] = 1 \forall n \geq 1$ satisfying $\psi \in [\underline{\psi}, \bar{\psi}]$ for $0 < \underline{\psi} \leq 1 \leq \bar{\psi} < \infty$ (and $\underline{\psi} = \bar{\psi} = 1$ is the degenerate case with no permanent shocks).

Following ?, in future periods $t+n \forall n \geq 1$ there is a small probability \wp that income will be zero (a ‘zero-income event’),

$$\xi_{t+n} = \begin{cases} 0 & \text{with probability } \wp > 0 \\ \theta_{t+n}/(1-\wp) & \text{with probability } (1-\wp) \end{cases} \quad (5)$$

where θ_{t+n} is an iid mean-one random variable ($\mathbb{E}_t[\theta_{t+n}] = 1 \forall n > 0$) whose distribution satisfies $\theta \in [\underline{\theta}, \bar{\theta}]$ where $0 < \underline{\theta} \leq 1 \leq \bar{\theta} < \infty$.⁸ Call the cumulative distribution functions \mathcal{F}_ψ and \mathcal{F}_θ (where \mathcal{F}_ξ is derived trivially from (5) and \mathcal{F}_θ). For quick identification in tables and graphs, we will call this the ‘Friedman/Muth’ model because it is a specific implementation of the ? model as interpreted by ?.

The model looks more special than it is. In particular, a positive probability of zero-income events may seem objectionable (despite empirical support).⁹ But a model with a nonzero minimum value of ξ (motivated, say, by the existence of unemployment insurance) can be redefined by capitalizing the PDV of minimum income into current market assets,¹⁰ transforming that model back into this one. And no key results would change if the transitory shocks were persistent but mean-reverting, instead of IID. Also, the assumption of a positive point mass for the worst realization of the transitory shock is inessential, but simplifies the proofs and is a powerful aid to intuition.

This model differs from Bewley’s (?) classic formulation in several ways. The CRRA utility function does not satisfy Bewley’s assumption that $u(0)$ is well-defined, or that $u'(0)$ is well defined and finite; indeed, neither the value function nor the marginal value function will be bounded. It differs from Schectman and Escudero (?) in that they impose liquidity constraints and positive minimum income. It differs from both of these in that it permits permanent growth in income, and also permanent shocks to income, which a large empirical literature finds to be of dominant importance in micro data.¹¹ It differs from Deaton (?) because liquidity constraints are absent; there are separate transitory and permanent shocks (*a la* ?); and the transitory shocks here can occasionally cause income to reach zero. It differs from models found in Stokey et. al. (?) because neither liquidity constraints nor bounds on utility or marginal utility are imposed.^{12,13} ?

⁸? and ? analyze cases where the shock processes have unbounded support.

⁹We calibrate this probability to 0.005 to match data from the Panel Study of Income Dynamics (?).

¹⁰So long as unemployment benefits are proportional to \mathbf{p}_t ; see the discussion in Section 2.11.

¹¹MaCurdy (?); Abowd and Card (?); Carroll and Samwick (?); Jappelli and Pistaferri (?). Much of the literature instead incorporates highly ‘persistent’ but not completely permanent shocks, but ? show that when measurement problems are handled correctly data yields serial correlation coefficients 0.98 – 1.00; and ? suggests that survey data support the same conclusion.

¹²Similar restrictions are made in the well known papers by Scheinkman and Weiss (?), Clarida (?), and ?. See ? for an elegant analysis of a related but simpler continuous-time model.

¹³? relaxed the bounds on the return function, but they address only the deterministic case.

show how to allow unbounded returns by using policy function iteration, but also impose constraints.

The paper with perhaps the most in common with this one is ?, henceforth MST, who establish the existence and uniqueness of a solution to a general income fluctuation problem in a Markovian setting. The most important differences are that MST impose liquidity constraints, assume that $u'(0) = 0$, and that expected marginal utility of income is finite ($\mathbb{E}[u'(Y)] < \infty$). These assumptions are not consistent with the combination of CRRA utility and income dynamics used here, whose joint properties are key to the results.¹⁴

2.2 The Problem Can Be Normalized By Permanent Income

We establish a bit more notation by reviewing the familiar result that in such problems (CRRA utility, permanent shocks) the number of states can be reduced from two (\mathbf{m} and \mathbf{p}) to one ($m = \mathbf{m}/\mathbf{p}$). Value in the last period is $u(\mathbf{m}_T)$; using (in the last line in (6) below) the fact that for our CRRA utility function, $u(xy) = x^{1-\rho}u(y)$, and generically defining nonbold variables as the boldface counterpart normalized by \mathbf{p}_t (as with $m = \mathbf{m}/\mathbf{p}$), consider the problem in the second-to-last period,

$$\begin{aligned} \mathbf{v}_{T-1}(\mathbf{m}_{T-1}, \mathbf{p}_{T-1}) &= \max_{\mathbf{c}_{T-1}} u(\mathbf{c}_{T-1}) + \beta \mathbb{E}_{T-1}[u(\mathbf{m}_T)] \\ &= \max_{c_{T-1}} u(\mathbf{p}_{T-1}c_{T-1}) + \beta \mathbb{E}_{T-1}[u(\mathbf{p}_T m_T)] \\ &= \mathbf{p}_{T-1}^{1-\rho} \left\{ \max_{c_{T-1}} u(c_{T-1}) + \beta \mathbb{E}_{T-1}[u(\Gamma_T m_T)] \right\}. \end{aligned} \quad (6)$$

Now, in a one-time deviation from the notational convention established in the last sentence, define nonbold ‘normalized value’ not as $\mathbf{v}_t/\mathbf{p}_t$ but as $v_t = \mathbf{v}_t/\mathbf{p}_t^{1-\rho}$, because this allows us to exploit features of the related problem,

$$\begin{aligned} v_t(m_t) &= \max_{\{c\}_t^T} u(c_t) + \beta \mathbb{E}_t[\Gamma_{t+1}^{1-\rho} v_{t+1}(m_{t+1})] \\ &\text{s.t.} \\ a_t &= m_t - c_t \\ k_{t+1} &= a_t/\Gamma_{t+1} \\ b_{t+1} &= k_{t+1}R = (R/\Gamma_{t+1})a_t = \mathcal{R}_{t+1}a_t \\ m_{t+1} &= b_{t+1} + \xi_{t+1}, \end{aligned} \quad (7)$$

where $\mathcal{R}_{t+1} \equiv (R/\Gamma_{t+1})$ is a ‘permanent-income-growth-normalized’ return factor, and the reformulated problem’s first order condition is¹⁵

¹⁴The incorporation of permanent shocks rules out application of the tools of ?, who followed and corrected an error in the fundamental work on the local contraction mapping method developed in ?. ? provide a correction to ?, that works under easier conditions to verify, but only addresses the deterministic case.

¹⁵Leaving aside their assumptions about the marginal utility function and liquidity constraints, it is tempting to view this as a special case of the model of MST, with our $\mathcal{R}_{t+1} = R/\Gamma_{t+1}$ (defined below equation (7)) corresponding to their stochastic rate of return on capital and the FVAF $\beta\Gamma_{t+1}^{1-\rho}$ defined below (39) corresponding to their stochastic discount factor. A caveat is that, here, \mathcal{R}_{t+1} and the modified discount factor are intimately (through Γ_{t+1}), which has profound

$$c_t^{-\rho} = R\beta \mathbb{E}_t[\Gamma_{t+1}^{-\rho} c_{t+1}^{-\rho}]. \quad (8)$$

Since $v_T(m_T) = u(m_T)$, defining $v_{T-1}(m_{T-1})$ from (7), we obtain

$$\mathbf{v}_{T-1}(\mathbf{m}_{T-1}, \mathbf{p}_{T-1}) = \mathbf{p}_{T-1}^{1-\rho} v_{T-1}(\underbrace{\mathbf{m}_{T-1}/\mathbf{p}_{T-1}}_{=m_{T-1}}).$$

This logic induces to earlier periods; if we solve the normalized one-state-variable problem (7), we will have solutions to the original problem for any $t < T$ from:

$$\begin{aligned} \mathbf{v}_t(\mathbf{m}_t, \mathbf{p}_t) &= \mathbf{p}_t^{1-\rho} v_t(m_t), \\ \mathbf{c}_t(\mathbf{m}_t, \mathbf{p}_t) &= \mathbf{p}_t c_t(m_t). \end{aligned}$$

2.3 Definition of a Nondegenerate Solution

The problem has a nondegenerate solution if as the horizon n gets arbitrarily large the solution in the first period of life $c_{T-n}(m)$ gets arbitrarily close to a limiting $c(m)$:

$$c(m) \equiv \lim_{n \rightarrow \infty} c_{T-n}(m) \quad (9)$$

that satisfies

$$0 < c(m) < \infty \quad (10)$$

for every $0 < m < \infty$.

2.4 Perfect Foresight Benchmarks

The familiar analytical solution to the perfect foresight model, obtained by setting $\wp = 0$ and $\underline{\theta} = \bar{\theta} = \underline{\psi} = \bar{\psi} = 1$, allows us to define some remaining notation and terminology.

2.4.1 Human Wealth

The dynamic budget constraint, strictly positive marginal utility, and the can't-die-in-debt condition (3) imply an exactly-holding intertemporal budget constraint (IBC):

$$\text{PDV}_t(\mathbf{c}) = \overbrace{\mathbf{m}_t - \mathbf{p}_t}^{\mathbf{b}_t} + \overbrace{\text{PDV}_t(\mathbf{p})}^{\mathbf{h}_t}, \quad (11)$$

where \mathbf{b} is nonhuman wealth, and with a constant $\mathcal{R} \equiv R/\Gamma$ ‘human wealth’ is

$$\begin{aligned} \mathbf{h}_t &= \mathbf{p}_t + \mathcal{R}^{-1} \mathbf{p}_t + \mathcal{R}^{-2} \mathbf{p}_t + \cdots + \mathcal{R}^{t-T} \mathbf{p}_t \\ &= \underbrace{\left(\frac{1 - \mathcal{R}^{-(T-t+1)}}{1 - \mathcal{R}^{-1}} \right)}_{\equiv h_t} \mathbf{p}_t. \end{aligned} \quad (12)$$

effects. It would be interesting, and should not be too difficult, to examine the case with independent shocks to the rate of return and productivity growth.

For $h \equiv \lim_{n \rightarrow \infty} h_{T-n}$ to be finite, need the Finite Human Wealth Condition (‘FHW’):

$$\underbrace{\Gamma/R}_{\equiv \mathcal{R}^{-1}} < 1. \quad (13)$$

Intuitively, finite human wealth requires a growth rate of (noncapital) income smaller than the interest rate at which that income is being discounted.

2.4.2 When Does the Perfect Foresight Unconstrained Solution Exist?

Without constraints, the consumption Euler equation always holds; with $u'(\mathbf{c}) = \mathbf{c}^{-\rho}$,

$$\mathbf{c}_{t+1}/\mathbf{c}_t = (\mathbf{R}\beta)^{1/\rho} \equiv \mathbf{P} \quad (14)$$

where the archaic letter ‘**thorn**’ represents what we will call the ‘Absolute Patience Factor’ or APF:

$$\mathbf{P} = (\mathbf{R}\beta)^{1/\rho}. \quad (15)$$

\mathbf{P} captures ‘patience’ because, if the ‘absolute impatience condition’ (AIC) holds,¹⁶

$$\mathbf{P} < 1, \quad (16)$$

the consumer’s level of spending will be too large to sustain indefinitely. We call such a consumer ‘absolutely impatient.’

A ‘Return Patience Factor’ (RPF) relates absolute patience to the return factor:

$$\mathbf{P}_R \equiv \mathbf{P}/\mathbf{R} \quad (17)$$

and since consumption is growing by \mathbf{P} but discounted by \mathbf{R} :

$$\text{PDV}_t(\mathbf{c}) = \left(\frac{1 - \mathbf{P}_R^{T-t+1}}{1 - \mathbf{P}_R} \right) \mathbf{c}_t \quad (18)$$

from which the IBC (11) implies

$$\mathbf{c}_t = \overbrace{\left(\frac{1 - \mathbf{P}_R}{1 - \mathbf{P}_R^{T-t+1}} \right)}^{\equiv \kappa_t} (\mathbf{b}_t + \mathbf{h}_t) \quad (19)$$

which defines a normalized finite-horizon perfect foresight consumption function

$$\bar{\mathbf{c}}_{T-n}(m_{T-n}) = \overbrace{(m_{T-n} - 1 + h_{T-n})}^{\equiv b_{T-n}} \underline{\kappa}_{T-n} \quad (20)$$

where $\underline{\kappa}_t$ is the marginal propensity to consume (MPC) — it answers the question ‘if the consumer had an extra unit of resources, how much more spending would occur.’ ($\bar{\mathbf{c}}$ ’s overbar signifies that $\bar{\mathbf{c}}$ will be an upper bound as we modify the problem to incorporate constraints and uncertainty; analogously, $\underline{\kappa}$ is a lower bound for the MPC).

¹⁶Impatience conditions have figured in intertemporal optimization problems since the beginning, e.g. in ?. These issues are so central that it would be hopeless to attempt to cite conditions in every other paper that correspond to conditions named and briefly exposted here. I make no claim to novelty for any condition aside from those implicated in my theorems, whose parallels *will* be articulated.

The denominator of (19) is the reason that, for $\underline{\kappa}$ to be strictly positive as $n = T - t$ goes to infinity, we must impose the Return Impatience Condition (RIC):

$$\mathbf{P}_R < 1, \quad (21)$$

so that

$$0 < \underline{\kappa} \equiv 1 - \mathbf{P}_R = \lim_{n \rightarrow \infty} \kappa_{T-n}. \quad (22)$$

The RIC thus implies that the consumer cannot be so pathologically patient as to wish, in the limit as the horizon approaches infinity, to spend nothing today out of an increase in current wealth (the RIC rules out the degenerate limiting solution $\bar{c}(m) = 0$). We call a consumer who satisfies the RIC ‘return impatient.’

Given that the RIC holds, and (as before) defining limiting objects by the absence of a time subscript, the limiting upper bound consumption function will be

$$\bar{c}(m) = (m + h - 1)\underline{\kappa}, \quad (23)$$

and so in order to rule out the degenerate limiting solution $\bar{c}(m) = \infty$ we need h to be finite; that is, we must impose the Finite Human Wealth Condition (13).

Because $u(xy) = x^{1-\rho}u(y)$ we can write a useful analytical expression for the value the consumer would achieve by spending permanent income \mathbf{p} in every period:

$$\begin{aligned} \mathbf{v}_t^{\text{autarky}} &= u(\mathbf{p}_t) + \beta u(\mathbf{p}_t \Gamma) + \beta^2 u(\mathbf{p}_t \Gamma^2) + \dots \\ &= u(\mathbf{p}_t) \left(1 + \beta \Gamma^{1-\rho} + (\beta \Gamma^{1-\rho})^2 + \dots \right) \\ &= u(\mathbf{p}_t) \left(\frac{1 - (\beta \Gamma^{1-\rho})^{T-t+1}}{1 - \beta \Gamma^{1-\rho}} \right) \end{aligned} \quad (24)$$

which (for $\Gamma > 0$) asymptotes to a finite number as $n = T - t$ approaches $+\infty$ if any of these equivalent conditions holds:

$$\begin{aligned} &\overbrace{\beta \Gamma^{1-\rho}}^{\equiv \beth} < 1 \\ &\beta R \Gamma^{-\rho} < R/\Gamma \\ &\mathbf{P}_R < (\Gamma/R)^{1-1/\rho}, \end{aligned} \quad (25)$$

where we call \beth^{17} the ‘Perfect Foresight Value Of Autarky Factor’ (PF-VAF), and the variants of (25) constitute alternative versions of the Perfect Foresight Finite Value of Autarky Condition, PF-FVAC; they guarantee that a consumer who always spends all permanent income ‘has finite autarky value’ (in the perfect foresight case).¹⁸

If the FHWC is satisfied, the PF-FVAC implies that the RIC is satisfied.¹⁹ Likewise,

¹⁷This is another kind of discount factor, so we use the Hebrew ‘bet’ which is a cognate of the Greek ‘beta’.

¹⁸This is related to the key impatience condition in ?.

¹⁹Divide both sides of the second inequality in (25) by R :

$$\mathbf{P}/R < (\Gamma/R)^{1-1/\rho} \quad (26)$$

and FHWC \Rightarrow the RHS is < 1 because $(\Gamma/R) < 1$ (and the RHS is raised to a positive power (because $\rho > 1$)).

if the FHWC and the GIC are both satisfied, PF-FVAC follows:

$$\begin{aligned}\mathbf{P} &< \Gamma < R \\ \mathbf{P}_R &< \Gamma/R < (\Gamma/R)^{1-1/\rho} < 1\end{aligned}\tag{27}$$

where the last line holds because $\text{FHWC} \Rightarrow 0 \leq (\Gamma/R) < 1$ and $\rho > 1 \Rightarrow 0 < 1-1/\rho < 1$.

The first panel of Table 4 summarizes: The PF-Unconstrained model has a non-degenerate limiting solution if we impose the RIC and FHWC (these conditions are necessary as well as sufficient). Together the PF-FVAC and the FHWC imply the RIC, so PF-FVAC and FHWC are jointly sufficient. If we impose the GIC and the FHWC, both the PF-FVAC and the RIC follow, so GIC+FHWC are also sufficient. But there are circumstances under which the RIC and FHWC can hold while the PF-FVAC fails (which we write ~~PF-FVAC~~). For example, if $\Gamma = 0$, the problem is a standard ‘cake-eating’ problem with a nondegenerate solution under the RIC.

Perhaps more useful than prose or a table, the relations of the conditions for the unconstrained perfect foresight case are presented diagrammatically in Figure 1. Each node represents a quantity considered in the foregoing analysis. The arrow associated with each inequality reflects the imposition of that condition. For example, one way we wrote the PF-FVAC in equation (25) is $\mathbf{P} < R^{1/\rho}\Gamma^{1-1/\rho}$, so imposition of the PF-FVAC is captured by the diagonal arrow connecting \mathbf{P} and $R^{1/\rho}\Gamma^{1-1/\rho}$. Traversing the boundary of the diagram clockwise starting at \mathbf{P} involves imposing first the GIC then the FHWC, and the consequent arrival at the bottom right node tells us that these two conditions jointly imply that the PF-FVAC holds. Reversal of a condition reverses the arrow’s direction; so, for example, the bottom-most arrow going from R to $R^{1/\rho}\Gamma^{1-1/\rho}$ imposes ~~FHWC~~; but we can cancel the cancellation and reverse the arrow. This would allow us to traverse the diagram in a clockwise direction from \mathbf{P} to R , revealing that imposition of GIC and FHWC (and, redundantly, FHWC again) let us conclude that the RIC holds because the starting point is \mathbf{P} and the endpoint is R . (Consult Appendix K for a detailed exposition of diagrams of this type).

2.4.3 PF Constrained Solution Exists Under RIC or Under~~{RIC,GIC}~~

We next examine the perfect foresight constrained solution because it is a useful benchmark (and limit) for the unconstrained problem with uncertainty (examined next).

If a liquidity constraint requiring $b \geq 0$ is ever to be relevant, it must be relevant at the lowest possible level of market resources, $m_t = 1$, defined by the lower bound for entering the period, $b_t = 0$ (if it were relevant at any higher point, it would certainly be relevant at this point). The constraint is ‘relevant’ if it prevents the choice that would otherwise be optimal; at $m_t = 1$ the constraint is relevant if the marginal utility from spending all of today’s resources $c_t = m_t = 1$, exceeds the marginal utility from doing the same thing next period, $c_{t+1} = 1$; that is, if such choices would violate the Euler equation (8):

$$1^{-\rho} > R\beta\Gamma^{-\rho}1^{-\rho}.\tag{28}$$



Figure 1 PF Unconstrained Model: Relation of GIC, FHWC, RIC, and PF-FVAC

An arrowhead points to the larger of the two quantities being compared. For example, the diagonal arrow indicates that $P < R^{1/\rho}\Gamma^{1-1/\rho}$, which is one way of writing the PF-FVAC, equation (25)

By analogy to the RPF, we therefore define a ‘growth patience factor’ (GPF) as

$$P_\Gamma = P/\Gamma, \quad (29)$$

and define a ‘growth impatience condition’ (GIC)

$$P_\Gamma < 1 \quad (30)$$

which is equivalent to (28) (exponentiate both sides by $1/\rho$).

We now examine implications of possible configurations of the conditions.

GIC and RIC. If the GIC fails but the RIC (21) holds, Appendix A shows that, for some $0 < m_\# < 1$, an unconstrained consumer behaving according to (23) would choose $c < m$ for all $m > m_\#$. In this case the solution to the constrained consumer’s problem is simple: For any $m \geq m_\#$ the constraint does not bind (and will never bind in the future); for such m the constrained consumption function is identical to the unconstrained one. If the consumer were somehow²⁰ to arrive at an $m < m_\# < 1$ the constraint would bind and the consumer would consume $c = m$. Using the $\dot{}$ accent for the version of a function \bullet in the presence of constraints (and recalling that $\bar{c}(m)$ is the unconstrained perfect foresight solution):

$$\dot{c}(m) = \begin{cases} m & \text{if } m < m_\# \\ \bar{c}(m) & \text{if } m \geq m_\#. \end{cases} \quad (31)$$

GIC and RIC. More useful is the case where the return impatience and GIC conditions

²⁰“Somehow” because $m < 1$ could only be obtained by entering the period with $b < 0$ which the constraint forbids.

both hold. In this case Appendix A shows that the limiting constrained consumption function is piecewise linear, with $\bar{c}(m) = m$ up to a first ‘kink point’ at $m_{\#}^1 > 1$, and with discrete declines in the MPC at a set of kink points $\{m_{\#}^1, m_{\#}^2, \dots\}$. As $m \uparrow \infty$ the constrained consumption function $\bar{c}(m)$ becomes arbitrarily close to the unconstrained $\bar{c}(m)$, and the marginal propensity to consume function $\bar{\kappa}(m) \equiv \bar{c}'(m)$ limits to $\underline{\kappa}$.²¹ Similarly, the value function $\bar{v}(m)$ is nondegenerate and limits into the value function of the unconstrained consumer.

This logic holds even when the finite human wealth condition fails ($\bar{\text{EHWC}}$), because the constraint prevents the (limiting) consumer²² from borrowing against unbounded human wealth to finance unbounded current consumption. Under these circumstances, the consumer who starts with any $b_t > 1$ will, over time, run those resources down so that by some finite number of periods τ in the future the consumer will reach $b_{t+\tau} = 0$, and thereafter will set $\mathbf{c} = \mathbf{p}$ for eternity (which the PF-FVAC says yields finite value). Using the same steps as for equation (24), value of the interim program is also finite:

$$\mathbf{v}_{t+\tau} = \Gamma^{\tau(1-\rho)} \mathbf{u}(\mathbf{p}_t) \left(\frac{1 - (\beta \Gamma^{1-\rho})^{T-(t+\tau)+1}}{1 - \beta \Gamma^{1-\rho}} \right).$$

So, if $\bar{\text{EHWC}}$, the limiting consumer’s value for any finite m will be the sum of two finite numbers: The component due to the unconstrained consumption choice made over the finite horizon leading up to $b_{t+\tau} = 0$, and the finite component due to the value of consuming all $\mathbf{p}_{t+\tau}$ thereafter.

GIC and $\bar{\text{RHC}}$. The most peculiar possibility occurs when the RIC fails. Under these circumstances the $\bar{\text{FWC}}$ must also fail (Appendix A), and the constrained consumption function is nondegenerate. (See appendix Figure 8 for a numerical example). While it is true that $\lim_{m \uparrow \infty} \bar{\kappa}(m) = 0$, nevertheless the limiting constrained consumption function $\bar{c}(m)$ is strictly positive and strictly increasing in m . This result interestingly reconciles the conflicting intuitions from the unconstrained case, where $\bar{\text{RHC}}$ would suggest a degenerate limit of $\bar{c}(m) = 0$ while $\bar{\text{EHWC}}$ would suggest a degenerate limit of $\bar{c}(m) = \infty$.

Tables 3 and 4 (and appendix table 5) codify.

We now examine the case with uncertainty but without constraints, which will turn out to be a close parallel to the model with constraints but without uncertainty.

2.5 Uncertainty-Modified Conditions

2.5.1 Impatience

When uncertainty is introduced, the expectation of beginning-of-period bank balances b_{t+1} can be rewritten as:

$$\mathbb{E}_t[b_{t+1}] = a_t \mathbb{E}_t[(\mathbf{R}/\Gamma_{t+1})] = a_t (\mathbf{R}/\Gamma) \mathbb{E}_t[\psi_{t+1}^{-1}] \quad (32)$$

²¹See ? for details.

²²That is, one obeying $c(m) = \lim_{n \uparrow \infty} c_{T-n}(m)$.

where Jensen's inequality guarantees that the expectation of the inverse of the permanent shock is greater than one. It will be convenient to define

$$\underline{\psi} \equiv (\mathbb{E}[\psi^{-1}])^{-1} \quad (33)$$

which satisfies $\underline{\psi} < 1$ (thanks again to Mr. Jensen), so we can define

$$\underline{\Gamma} \equiv \Gamma \underline{\psi} < \Gamma \quad (34)$$

which allows us to write uncertainty-adjusted versions of equations and conditions in a manner exactly parallel to those for the perfect foresight case; for example, we define a normalized Growth Patience Factor (GPF-Nrm):

$$\mathbf{P}_{\underline{\Gamma}} = \mathbf{P}/\underline{\Gamma} = \mathbb{E}[\mathbf{P}/(\Gamma\psi)] \quad (35)$$

and a normalized version of the Growth Impatience Condition, GIC-Nrm:

$$\mathbf{P}_{\underline{\Gamma}} < 1, \quad (36)$$

that is stronger than the perfect foresight version (30) because $\underline{\Gamma} < \Gamma$.

2.5.2 Autarky Value

Analogously to (24), value for a consumer who spent exactly their permanent income every period would reflect the product of the expectation of the (independent) future shocks to permanent income:

$$\begin{aligned} \mathbf{v}_t &= \mathbb{E}_t [\mathbf{u}(\mathbf{p}_t) + \beta \mathbf{u}(\mathbf{p}_t \Gamma_{t+1}) + \cdots + \beta^{T-t} \mathbf{u}(\mathbf{p}_t \Gamma_{t+1} \cdots \Gamma_T)] \\ &= \mathbf{u}(\mathbf{p}_t) \left(\frac{1 - (\beta \Gamma^{1-\rho} \mathbb{E}[\psi^{1-\rho}])^{T-t+1}}{1 - \beta \Gamma^{1-\rho} \mathbb{E}[\psi^{1-\rho}]} \right), \end{aligned}$$

suggesting the definition of a utility-compensated equivalent of the permanent shock,

$$\underline{\underline{\psi}} = (\mathbb{E}[\psi^{1-\rho}])^{1/(1-\rho)}, \quad (37)$$

which will satisfy $\underline{\underline{\psi}} < 1$ for $\rho > 1$ and nondegenerate ψ . Defining

$$\underline{\underline{\Gamma}} = \Gamma \underline{\underline{\psi}} \quad (38)$$

we can see that \mathbf{v}_t will be positive and finite as T approaches ∞ if

$$\begin{aligned} 0 &< \overbrace{\beta \underline{\underline{\Gamma}}^{1-\rho}}^{\underline{\underline{\Xi}}} < 1 \\ 0 &< \beta < \underline{\underline{\Gamma}}^{\rho-1} \end{aligned} \quad (39)$$

which we call the ‘finite value of autarky condition’ because it guarantees that value is finite for a consumer who always consumes their (now stochastic) permanent income (and we will call $\underline{\underline{\Xi}}$ the ‘Value of Autarky Factor’ (or ‘VAF’)).²³ For nondegenerate ψ ,

²³In a stationary environment — that is, with $\underline{\underline{\Gamma}} = 1$ — this corresponds to an impatience condition imposed by ?; but their remaining conditions do not correspond to those here, because their problem differs and their method of proof differs.

Table 1 Microeconomic Model Calibration

Calibrated Parameters			
Description	Parameter	Value	Source
Permanent Income Growth Factor	Γ	1.03	PSID: Carroll (1992)
Interest Factor	R	1.04	Conventional
Time Preference Factor	β	0.96	Conventional
Coefficient of Relative Risk Aversion	ρ	2	Conventional
Probability of Zero Income	\wp	0.005	PSID: Carroll (1992)
Std Dev of Log Permanent Shock	σ_ψ	0.1	PSID: Carroll (1992)
Std Dev of Log Transitory Shock	σ_θ	0.1	PSID: Carroll (1992)

this condition is stronger (harder to satisfy in the sense of requiring lower β) than the perfect foresight version (25) because $\underline{\Gamma} < \Gamma$.²⁴

2.6 The Baseline Numerical Solution

Figure 2, familiar from the literature, depicts the successive consumption rules that apply in the last period of life ($c_T(m)$), the second-to-last period, and earlier periods under baseline parameter values listed in Table 2. (The 45 degree line is $c_T(m) = m$ because in the last period of life it is optimal to spend all remaining resources.)

In the figure, the consumption rules appear to converge to a nondegenerate $c(m)$. Our next purpose is to show that this appearance is not deceptive.

2.7 Concave Consumption Function Characteristics

A precondition for the main proof is that the maximization problem defines a sequence of continuously differentiable strictly increasing strictly concave²⁵ functions $\{c_T, c_{T-1}, \dots\}$. The straightforward but tedious proof is relegated to Appendix B. For present purposes, the most important point is that the income process induces what ? dubbed a ‘natural borrowing constraint’: $c_t(m) < m$ for all periods $t < T$ because a consumer who spent all available resources would arrive in period $t+1$ with balances b_{t+1} of zero, and then might earn zero income over the remaining horizon, risking the possibility of a requirement to spend zero, risking negative infinite utility. To avoid this disaster, the consumer never

²⁴To see this, rewrite (39) as

$$\begin{aligned}
\beta R &< R \underline{\Gamma}^{\rho-1} \\
(\beta R)^{1/\rho} &< R^{1/\rho} \underline{\Gamma}^{1-1/\rho} \underline{\psi}^{1-1/\rho} \\
\mathbf{P}_\Gamma &< (R/\Gamma)^{1/\rho} \underline{\psi}^{1-1/\rho}
\end{aligned}$$

where the last equation is the same as the PF-FVAC condition except that the RHS is multiplied by $\underline{\psi}^{1-1/\rho}$ which is strictly less than 1.

²⁵With one obvious exception: $c_T(m)$ is linear (and so only weakly concave).

Table 2 Model Characteristics Calculated from Parameters

Description	Symbol and Formula	Approximate Calculated Value
Finite Human Wealth Factor	$\mathcal{R}^{-1} \equiv \Gamma/R$	0.990
PF Finite Value of Autarky Factor	$\sqsupset \equiv \beta\Gamma^{1-\rho}$	0.932
Growth Compensated Permanent Shock	$\underline{\psi} \equiv (\mathbb{E}[\psi^{-1}])^{-1}$	0.990
Uncertainty-Adjusted Growth	$\underline{\Gamma} \equiv \Gamma\underline{\psi}$	1.020
Utility Compensated Permanent Shock	$\underline{\underline{\psi}} \equiv (\mathbb{E}[\psi^{1-\rho}])^{1/(1-\rho)}$	0.990
Utility Compensated Growth	$\underline{\underline{\Gamma}} \equiv \Gamma\underline{\underline{\psi}}$	1.020
Absolute Patience Factor	$\mathfrak{P} \equiv (R\beta)^{1/\rho}$	0.999
Return Patience Factor	$\mathfrak{P}_R \equiv \mathfrak{P}/R$	0.961
Growth Patience Factor	$\mathfrak{P}_\Gamma \equiv \mathfrak{P}/\Gamma$	0.970
Normalized Growth Patience Factor	$\mathfrak{P}_{\underline{\Gamma}} \equiv \mathfrak{P}/\underline{\Gamma}$	0.980
Finite Value of Autarky Factor	$\sqsubseteq \equiv \beta\Gamma^{1-\rho}\underline{\underline{\psi}}^{1-\rho}$	0.941
Weak Impatience Factor	$\wp^{1/\rho}\mathfrak{P} \equiv (\wp\beta R)^{1/\rho}$	0.071



Figure 2 Convergence of the Consumption Rules

spends everything. ? seems to have been the first to argue, based on his numerical results, that the natural borrowing constraint was a quantitatively plausible alternative to ‘artificial’ or ‘ad hoc’ borrowing constraints.²⁶

Strict concavity and continuous differentiability of the consumption function are key elements in many of the arguments below, but are not characteristics of models with ‘artificial’ borrowing constraints – and we will see below that the analytical convenience of these features is considerable.

2.8 Bounds for the Consumption Functions

The consumption functions depicted in Figure 2 appear to have limiting slopes as $m \downarrow 0$ and as $m \uparrow \infty$. This section confirms that impression and derives those slopes, which will be needed in the contraction mapping proof.²⁷

Assume (as justified above) that a continuously differentiable concave consumption function exists in period $t + 1$, with an origin at $c_{t+1}(0) = 0$, a minimal MPC $\underline{\kappa}_{t+1} > 0$, and maximal MPC $\bar{\kappa}_{t+1} \leq 1$. (If $t + 1 = T$ these will be $\underline{\kappa}_T = \bar{\kappa}_T = 1$; for earlier periods they will exist by recursion.)

Under our imposed assumption that human wealth is finite, the MPC bound as wealth approaches infinity is easy to understand: As the *proportion* of consumption that will be financed out of human wealth approaches zero, the proportional difference between the solution to the model with uncertainty and the perfect foresight model shrinks to zero. In the course of proving this, Appendix G provides a useful recursive expression (used below) for the (inverse of the) limiting MPC:

$$\underline{\kappa}_t^{-1} = 1 + \mathbf{P}_R \underline{\kappa}_{t+1}^{-1}. \quad (40)$$

2.8.1 Weak RIC Conditions

Equation (85) presents a parallel expression for the limiting maximal MPC as $m_t \downarrow 0$:

$$\bar{\kappa}_t^{-1} = 1 + \wp^{1/\rho} \mathbf{P}_R \bar{\kappa}_{t+1}^{-1} \quad (41)$$

where $\{\bar{\kappa}_{T-n}^{-1}\}_{n=0}^{\infty}$ is a decreasing convergent sequence if the ‘weak return patience factor’ $\wp^{1/\rho} \mathbf{P}_R$ satisfies:

$$0 \leq \wp^{1/\rho} \mathbf{P}_R < 1, \quad (42)$$

a condition we dub the ‘Weak Return Impatience Condition’ (WRIC) because with $\wp < 1$ it will hold more easily (for a larger set of parameter values) than the RIC ($\mathbf{P}_R < 1$). The essence of the argument is that as wealth approaches zero, the overriding consideration

²⁶The same (numerical) point applies for infinite horizon models (calibrated to actual empirical data on household income dynamics); cf. ?.

²⁷? show that the consumption function becomes linear as wealth approaches infinity in a model with capital income risk and liquidity constraints; ? show that these results generalize to the limits derived here if capital income is added to the model.

that limits consumption is the (recursive) fear of the zero-income events. (That is why the probability of the zero income event \wp appears in the expression.)

We are now in position to observe that the optimal consumption function must satisfy

$$\underline{\kappa}_t m_t \leq c_t(m_t) \leq \bar{\kappa}_t m_t \quad (43)$$

because consumption starts at zero and is continuously differentiable, is strictly concave,²⁸ and always exhibits a slope between $\underline{\kappa}_t$ and $\bar{\kappa}_t$ (the formal proof is in Appendix D).

2.9 Conditions Under Which the Problem Defines a Contraction Mapping

As mentioned above, standard theorems in the contraction mapping literature following Stokey et. al. (?) require utility or marginal utility to be bounded over the space of possible values of m , which does not hold here because the possibility (however unlikely) of an unbroken string of zero-income events through the end of the horizon means that utility (and marginal utility) are unbounded as $m \downarrow 0$. Although a recent literature examines the existence and uniqueness of solutions to Bellman equations in the presence of ‘unbounded returns’ (see, e.g., ?), the techniques in that literature cannot be used to solve the problem here because the required conditions are violated by a problem that incorporates permanent shocks.²⁹

Fortunately, Boyd (?) provided a weighted contraction mapping theorem that ? showed could be used to address the homogeneous case (of which CRRA is an example) in a deterministic framework; later, ? showed how to extend the ? approach to the stochastic case.

Definition 1. Consider any function $\bullet \in \mathcal{C}(\mathcal{A}, \mathcal{B})$ where $\mathcal{C}(\mathcal{A}, \mathcal{B})$ is the space of continuous functions from \mathcal{A} to \mathcal{B} . Suppose $F \in \mathcal{C}(\mathcal{A}, \mathcal{B})$ with $\mathcal{B} \subseteq \mathbb{R}$ and $F > 0$. Then \bullet is F -bounded if the F -norm of \bullet ,

$$\|\bullet\|_F = \sup_m \left[\frac{|\bullet(m)|}{F(m)} \right], \quad (44)$$

is finite.

For $\mathcal{C}_F(\mathcal{A}, \mathcal{B})$ defined as the set of functions in $\mathcal{C}(\mathcal{A}, \mathcal{B})$ that are F -bounded; w, x, y, and z as examples of F -bounded functions; and using $\mathbf{0}(m) = 0$ to indicate the function that returns zero for any argument, Boyd (?) proves the following.

Boyd’s Weighted Contraction Mapping Theorem. Let $T : \mathcal{C}_F(\mathcal{A}, \mathcal{B}) \rightarrow \mathcal{C}(\mathcal{A}, \mathcal{B})$ such that^{30,31}

- (1) T is non-decreasing, i.e. $x \leq y \Rightarrow \{Tx\} \leq \{Ty\}$
- (2) $\{T\mathbf{0}\} \in \mathcal{C}_F(\mathcal{A}, \mathcal{B})$

²⁸?

²⁹See ? for a detailed discussion of the reasons the existing literature up through ? cannot handle the problem described here.

³⁰We will usually denote the function that results from the mapping as, e.g., $\{Tw\}$.

³¹To non-theorists, this notation may be slightly confusing; the inequality relations in (1) and (3) are taken to mean ‘for any specific element \bullet in the domain of the functions in question’ so that, e.g., $x \leq y$ is short for $x(\bullet) \leq y(\bullet) \forall \bullet \in \mathcal{A}$. In this notation, $\zeta \alpha F$ in (3) is a *function* which can be applied to any argument \bullet (because F is a function).

(3) *There exists some real $0 < \alpha < 1$ such that*

$$\{\mathsf{T}(\mathbf{w} + \zeta F)\} \leq \{\mathsf{T}\mathbf{w}\} + \zeta \alpha F \quad \text{holds for all real } \zeta > 0 \quad .$$

Then T defines a contraction with a unique fixed point.

For our problem, take \mathcal{A} as $\mathbb{R}_{>0}$ and \mathcal{B} as \mathbb{R} , and define

$$\{\mathsf{Ez}\}(a_t) = \mathbb{E}_t \left[\Gamma_{t+1}^{1-\rho} z(a_t \mathcal{R}_{t+1} + \xi_{t+1}) \right] .$$

Using this, we introduce the mapping $\mathcal{T} : \mathcal{C}_F(\mathcal{A}, \mathcal{B}) \rightarrow \mathcal{C}(\mathcal{A}, \mathcal{B})$,³²

We can show that our operator \mathcal{T} satisfies the conditions that Boyd requires of his operator T , if we impose two restrictions on parameter values. The first is the WRIC necessary for convergence of the maximal MPC, equation (42) above. More serious is the Finite Value of Autarky condition, equation (39). (We discuss the interpretation of these restrictions in detail in Section 2.11 below.) Imposing these restrictions, we are now in position to state the central theorem of the paper.

Theorem 1. *\mathcal{T} is a contraction mapping if the restrictions on parameter values (42) and (39) are true (that is, if the weak return impatience condition and the finite value of autarky condition hold).*

Intuitively, Boyd’s theorem shows that if you can find a F that is everywhere finite but goes to infinity ‘as fast or faster’ than the function you are normalizing with F , the normalized problem defines a contraction mapping. The intuition for the FVAC condition is that, with an infinite horizon, with any initial amount of bank balances b_0 , in the limit your value can always be made greater than you would get by consuming exactly the sustainable amount (say, by consuming $(r/R)b_0 - \epsilon$ for some small $\epsilon > 0$).

The cumbersome details of the proof are relegated to Appendix D. Given that the value function converges, Appendix E.2 shows that the consumption functions converge.³³

2.10 The Liquidity Constrained Solution as a Limit

This section explains why a related problem commonly considered in the literature (e.g., by Deaton (?)), with a liquidity constraint and a positive minimum value of income, is the limit of the problem considered here as the probability \wp of the zero-income event approaches zero.

The ‘related’ problem makes two changes to the problem defined above:

1. An ‘artificial’ liquidity constraint is imposed: $a_t \geq 0$
2. The probability of zero-income events is zero: $\wp = 0$

The essence of the argument is simple. Imposing the artificial constraint without changing \wp would not change behavior at all: The possibility of earning zero income

³²Note that the existence of the maximum is assured by the continuity of $\{\mathsf{Ez}\}(a_t)$ (it is continuous because it is the sum of continuous F -bounded functions z) and the compactness of $[\underline{\kappa}m_t, \bar{\kappa}m_t]$.

³³MST’s proof is for convergence of the consumption policy function directly, rather than of the value function, which is why their conditions are on u' , which governs behavior.

over the remaining horizon already prevents the consumer from ending the period with zero assets. So, for precautionary reasons, the consumer will save something.

But the *extent* to which the consumer feels the need to make this precautionary provision depends on the *probability* that it will turn out to matter. As $\wp \downarrow 0$, that probability becomes arbitrarily small, so the *amount* of precautionary saving induced by the zero-income events approaches zero as $\wp \downarrow 0$. But “zero” is the amount of precautionary saving that would be induced by a zero-probability event for the impatient liquidity constrained consumer.

Another way to understand this is just to think of the liquidity constraint reflecting a component of the utility function that is zero whenever the consumer ends the period with (strictly) positive assets, but negative infinity if the consumer ends the period with (weakly) negative assets.

See Appendix H for the formal proof justifying the foregoing intuitive discussion.³⁴

The conditions required for convergence and nondegeneracy are thus strikingly similar between the liquidity constrained perfect foresight model and the model with uncertainty but no explicit constraints: The liquidity constrained perfect foresight model is just the limiting case of the model with uncertainty as the degree of all three kinds of uncertainty (zero-income events, other transitory shocks, and permanent shocks) approaches zero.

2.11 Discussion of Parametric Restrictions

The full relationship among conditions is represented in Figure 3. Though the diagram looks complex, it is merely a modified version of the earlier diagram (Figure 1) with further (mostly intermediate) inequalities inserted. (Arrows with a “because” now label relations that always hold under the model’s assumptions.)³⁵

2.11.1 The WRIC

The ‘weakness’ of the additional condition sufficient for contraction beyond the FVAC, the WRIC, can be seen by asking ‘under what circumstances would the FVAC hold but the WRIC fail?’ Algebraically, the requirement is

$$\beta \Gamma^{1-\rho} \underline{\underline{\psi}}^{1-\rho} < 1 < (\wp \beta)^{1/\rho} / R^{1-1/\rho}. \quad (45)$$

If we require $R \geq 1$, the WRIC is redundant because now $\beta < 1 < R^{\rho-1}$, so that (with $\rho > 1$ and $\beta < 1$) the RIC (and WRIC) must hold. But neither theory nor evidence demand that $R \geq 1$. We can therefore approach the question of the WRIC’s relevance by asking just how low R must be for the condition to be relevant. Suppose for illustration that $\rho = 2$, $\underline{\underline{\psi}}^{1-\rho} = 1.01$, $\Gamma^{1-\rho} = 1.01^{-1}$ and $\wp = 0.10$. In that case (45) reduces to

$$\beta < 1 < (0.1\beta/R)^{1/2}$$

³⁴It seems likely that a similar argument would apply even in the context of a model like that of MST, perhaps with some weak restrictions on returns.

³⁵Again, readers unfamiliar with such diagrams should see Appendix K for a more detailed exposition.



Figure 3 Relation of All Inequality Conditions

See Table 2 for Numerical Values of Nodes Under Baseline Parameters

but since $\beta < 1$ by assumption, the binding requirement is that

$$R < \beta/10$$

so that for example if $\beta = 0.96$ we would need $R < 0.096$ (that is, a perpetual riskfree rate of return of worse than -90 percent a year) in order for the WRIC to bind.

Perhaps the best way of thinking about this is to note that the space of parameter values for which the WRIC is relevant shrinks out of existence as $\varphi \rightarrow 0$, which Section 2.10 showed was the precise limiting condition under which behavior becomes arbitrarily close to the liquidity constrained solution (in the absence of other risks). On the other hand, when $\varphi = 1$, the consumer has no noncapital income (so that the FHWC holds) and with $\varphi = 1$ the WRIC is identical to the RIC; but the RIC is the only condition required for a solution to exist for a perfect foresight consumer with no noncapital income. Thus the WRIC forms a sort of ‘bridge’ between the liquidity constrained and the unconstrained problems as φ moves from 0 to 1.

2.11.2 When the RIC Fails

In the perfect foresight problem (Section 2.4.2), the RIC was necessary for existence of a nondegenerate solution. It is surprising, therefore, that in the presence of uncertainty, the much weaker WRIC is sufficient for nondegeneracy (assuming that the FVAC holds). We can directly derive the features the problem must exhibit (given the FVAC) under

~~RIC~~ (that is, $R < (R\beta)^{1/\rho}$):

$$\begin{aligned}
& \overbrace{R < (R\beta)^{1/\rho} < (R(\Gamma\underline{\psi})^{\rho-1})^{1/\rho}}^{\text{implied by FVAC}} \\
& R < (R/\Gamma)^{1/\rho} \Gamma \underline{\psi}^{1-1/\rho} \\
& R/\Gamma < (R/\Gamma)^{1/\rho} \underline{\psi}^{1-1/\rho} \\
& R/\Gamma < \underline{\underline{\psi}}
\end{aligned} \tag{46}$$

but since $\underline{\underline{\psi}} < 1$ (cf. the argument below (37)), this requires $R/\Gamma < 1$; so, given the FVAC, the RIC can fail only if human wealth is unbounded. As an illustration of the usefulness of our diagrams, note that this algebraically complicated conclusion could be easily reached diagrammatically in figure 3 by starting at the R node and imposing ~~RIC~~, reversing the RIC arrow and then traversing the diagram along any clockwise path to the PF-VAF node at which point we realize that we *cannot* impose the FHWC because that would let us conclude $R > R$.

As in the perfect foresight constrained problem, unbounded limiting human wealth (~~FHWC~~) here does not lead to a degenerate limiting consumption function (finite human wealth is not a condition required for the convergence theorem). But, from equation (40) and the discussion surrounding it, an implication of ~~RIC~~ is that $\lim_{m \uparrow \infty} c'(m) = 0$. Thus, interestingly, in the special $\{\text{RIC}, \text{FHWC}\}$ case (unavailable in the perfect foresight model) the presence of uncertainty both permits unlimited human wealth (in the $n \uparrow \infty$ limit) and at the same time prevents unlimited human wealth from resulting in (limiting) infinite consumption at any finite m . Intuitively, in the presence of uncertainty, pathological patience (which in the perfect foresight model results in a limiting consumption function of $c(m) = 0$ for finite m) plus unbounded human wealth (which the perfect foresight model prohibits because it leads to a limiting consumption function $c(m) = \infty$ for any finite m) combine to yield a unique finite level of consumption and MPC for any finite value of m . Note the close parallel to the conclusion in the perfect foresight liquidity constrained model in the $\{\text{GIC}, \text{RIC}\}$ case. There, too, the tension between infinite human wealth and pathological patience was resolved with a nondegenerate consumption function whose limiting MPC was zero.³⁶

2.11.3 When the RIC Holds

FHWC. If the RIC and FHWC both hold, a perfect foresight solution exists (see 2.4.2 above). As $m \uparrow \infty$ the limiting consumption function and value function become arbitrarily close to those in the perfect foresight model, because human wealth pays for a vanishingly small portion of spending. This will be the main case analyzed in detail below.

FHWC. The more exotic case is where FHWC does not hold; in the perfect foresight model, $\{\text{RIC}, \text{FHWC}\}$ is the degenerate case with limiting $\bar{c}(m) = \infty$. Here, since the

³⁶? derive conditions under which the limiting MPC is zero in an even more general case where there is also capital income risk.

FVAC implies that the PF-FVAC holds (traverse Figure 3 clockwise from \mathbf{P} by imposing FVAC and continue to the PF-VAF node): Reversing the arrow connecting the R and PF-VAF nodes implies that under \mathbf{EHWC} :

$$\begin{array}{c} \text{PF-FVAC} \\ \hline \mathbf{P} < (\mathbf{R}/\Gamma)^{1/\rho} \Gamma \\ \mathbf{P} < \Gamma \end{array}$$

where the transition from the first to the second lines is justified because $\mathbf{EHWC} \Rightarrow (\mathbf{R}/\Gamma)^{1/\rho} < 1$. So, $\{\mathbf{RIC}, \mathbf{EHWC}\}$ implies the GIC holds. However, we are not entitled to conclude that the GIC-Nrm holds: $\mathbf{P} < \Gamma$ does not imply $\mathbf{P} < \psi\Gamma$ where $\psi < 1$.

We have now established the principal points of comparison between the perfect foresight solutions and the solutions under uncertainty; these are codified in the remaining parts of Tables 3 and 4.

3 Analysis of the Converged Consumption Function

Figures 4-6 capture the main properties of the converged consumption rule when the RIC, GIC-Nrm, and FHWC all hold.³⁷

Figure 4 shows the expected growth factors for consumption, the level of market resources, and the market resources ratio, $\mathbb{E}_t[\mathbf{c}_{t+1}/\mathbf{c}_t]$ and $\mathbb{E}_t[\mathbf{m}_{t+1}/\mathbf{m}_t]$, and $\mathbb{E}_t[m_{t+1}/m_t]$, for a consumer behaving according to the converged consumption rule, while Figures 5—6 illustrate theoretical bounds for the consumption function and the MPC.

Three features of behavior are captured, or suggested, by the figures.

First, as $m_t \uparrow \infty$ the expected consumption growth factor goes to \mathbf{P} , indicated by the lower bound in Figure 4, and the marginal propensity to consume approaches $\underline{\kappa} = (1 - \mathbf{P}_R)$ (see Figure 5) — the same as the perfect foresight MPC. Second, as m_t approaches zero the consumption growth factor approaches ∞ (Figure 4) and the MPC approaches $\bar{\kappa} = (1 - \wp^{1/\rho}\mathbf{P}_R)$ (Figure 5). Third, there is a value $m = \hat{m}$ at which the expected growth rate of \mathbf{m} matches the expected growth rate of permanent income Γ , and a different (lower) value where the expected growth rate of consumption at $m = \hat{m}$ is lower than Γ . Thus, at the individual level, this model does not have a ‘balanced growth’ equilibrium.³⁸

3.1 Limits as m approaches Infinity

Define

$$\underline{c}(m) = \underline{\kappa}m$$

³⁷These figures reflect the converged rule corresponding to the parameter values indicated in Table 2.

³⁸A final proposition suggested by Figure 4 is that the expected consumption growth factor is declining in the level of the cash-on-hand ratio m_t . This turns out to be true in the absence of permanent shocks, but in extreme cases it can be false if permanent shocks are present; see Appendix L.

Table 3 Definitions and Comparisons of Conditions

Perfect Foresight Versions	Uncertainty Versions
Finite Human Wealth Condition (FHC)	
$\Gamma/R < 1$ The growth factor for permanent income Γ must be smaller than the discounting factor R for human wealth to be finite.	$\Gamma/R < 1$ The model's risks are mean-preserving spreads, so the PDV of future income is unchanged by their introduction.
Absolute Impatience Condition (AIC)	
$\mathbf{P} < 1$ The unconstrained consumer is sufficiently impatient that the level of consumption will be declining over time: $\mathbf{c}_{t+1} < \mathbf{c}_t$	$\mathbf{P} < 1$ <i>If wealth is large enough, the expectation of consumption next period will be smaller than this period's consumption:</i> $\lim_{m_t \rightarrow \infty} \mathbb{E}_t[\mathbf{c}_{t+1}] < \mathbf{c}_t$
Return Impatience Conditions	
Return Impatience Condition (RIC)	Weak RIC (WRIC)
$\mathbf{P}/R < 1$ The growth factor for consumption \mathbf{P} must be smaller than the discounting factor R , so that the PDV of current and future consumption will be finite: $c'(m) = 1 - \mathbf{P}/R < 1$	$\wp^{1/\rho} \mathbf{P}/R < 1$ If the probability of the zero-income event is $\wp = 1$ then income is always zero and the condition becomes identical to the RIC. Otherwise, weaker. $c'(m) < 1 - \wp^{1/\rho} \mathbf{P}/R < 1$
Growth Impatience Conditions	
GIC	GIC-Nrm
$\mathbf{P}/\Gamma < 1$ For an unconstrained PF consumer, the ratio of \mathbf{c} to \mathbf{p} will fall over time. For constrained, guarantees the constraint eventually binds. Guarantees $\lim_{m_t \uparrow \infty} \mathbb{E}_t[\psi_{t+1} m_{t+1}/m_t] = \mathbf{P}_\Gamma$	$\mathbf{P} \mathbb{E}[\psi^{-1}]/\Gamma < 1$ By Jensen's inequality stronger than GIC. Ensures consumers will not expect to accumulate m unboundedly. $\lim_{m_t \rightarrow \infty} \mathbb{E}_t[m_{t+1}/m_t] = \mathbf{P}_\Gamma$
Finite Value of Autarky Conditions	
PF-FVAC	FVAC
$\beta \Gamma^{1-\rho} < 1$ equivalently $\mathbf{P} < R^{1/\rho} \Gamma^{1-1/\rho}$ The discounted utility of constrained consumers who spend their permanent income each period should be finite.	$\beta \Gamma^{1-\rho} \mathbb{E}[\psi^{1-\rho}] < 1$ By Jensen's inequality, stronger than the PF-FVAC because for $\rho > 1$ and nondegenerate ψ , $\mathbb{E}[\psi^{1-\rho}] > 1$.

Table 4 Sufficient Conditions for Nondegenerate[‡] Solution

Consumption Model(s)	Conditions	Comments
$\bar{c}(m)$: PF Unconstrained $\underline{c}(m) = \underline{\kappa}m$ Section 2.4.2: Section 2.4.2: Eq (26): Eq (27):	RIC, FHW ^o	RIC $\Rightarrow v(m) < \infty$; FHW ^o $\Rightarrow 0 < v(m) $ PF model with no human wealth ($h = 0$) RIC prevents $\bar{c}(m) = \underline{c}(m) = 0$ FHW ^o prevents $\bar{c}(m) = \infty$ PF-FVAC+FHW ^o \Rightarrow RIC GIC+FHW ^o \Rightarrow PF-FVAC
$\dot{c}(m)$: PF Constrained Section 2.4.3: Appendix A: Appendix A:	GIC , RIC GIC, RIC GIC, RIC	FHW ^o holds ($\Gamma < \mathbf{P} < R \Rightarrow \Gamma < R$) $\dot{c}(m) = \bar{c}(m)$ for $m > m_{\#} < 1$ (RIC would yield $m_{\#} = 0$ so $\dot{c}(m) = 0$) $\lim_{m \rightarrow \infty} \dot{c}(m) = \bar{c}(m)$, $\lim_{m \rightarrow \infty} \dot{\kappa}(m) = \underline{\kappa}$ kinks where horizon to $b = 0$ changes* $\lim_{m \rightarrow \infty} \dot{\kappa}(m) = 0$ kinks where horizon to $b = 0$ changes*
$c(m)$: Friedman/Muth Section 2.9: Section 2.11.1: Figure 3: Section 2.11.3: Section 2.11.2: Section 3.3: Section 3.3.2: Section 3.3.1:	Section 3.1, Section 3.2 FVAC, WRIC	$\underline{c}(m) < c(m) < \bar{c}(m)$ $\underline{v}(m) < v(m) < \bar{v}(m)$ Sufficient for Contraction WRIC is weaker than RIC FVAC is stronger than PF-FVAC FHW^o +RIC \Rightarrow GIC, $\lim_{m \rightarrow \infty} \kappa(m) = \underline{\kappa}$ RIC \Rightarrow FHW^o , $\lim_{m \rightarrow \infty} \kappa(m) = 0$ “Buffer Stock Saving” Conditions GIC $\Rightarrow \exists 0 < \tilde{m} < \infty$ GIC-Nrm $\Rightarrow \exists 0 < \hat{m} < \infty$

[‡]For feasible m satisfying $0 < m < \infty$, a nondegenerate limiting consumption function defines a unique optimal value of c satisfying $0 < c(m) < \infty$; a nondegenerate limiting value function defines a corresponding unique value of $-\infty < v(m) < 0$. ^oRIC, FHW^o are necessary as well as sufficient for the perfect foresight case. *That is, the first kink point in $c(m)$ is $m_{\#}$ s.t. for $m < m_{\#}$ the constraint will bind now, while for $m > m_{\#}$ the constraint will bind one period in the future. The second kink point corresponds to the m where the constraint will bind two periods in the future, etc. **In the Friedman/Muth model, the RIC+FHW^o are sufficient, but *not* necessary for nondegeneracy




Figure 4 ‘Stable’ m Values and Expected Growth Factors

which is the solution to an infinite-horizon problem with no noncapital income ($\xi_{t+n} = 0 \forall n \geq 1$); clearly $\underline{c}(m) < c(m)$, since allowing the possibility of future noncapital income cannot reduce current consumption. Our imposition of the RIC guarantees that $\underline{\kappa} > 0$, so this solution satisfies our definition of nondegeneracy, and because this solution is always available it defines a lower bound on both the consumption and value functions.

Assuming the FHCW holds, the infinite horizon perfect foresight solution (23) constitutes an upper bound on consumption in the presence of uncertainty, since the introduction of uncertainty strictly decreases the level of consumption at any m (?). Thus, we can write

$$\begin{aligned} \underline{c}(m) &< c(m) < \bar{c}(m) \\ 1 &< c(m)/\underline{c}(m) < \bar{c}(m)/\underline{c}(m). \end{aligned} \tag{47}$$

But

$$\begin{aligned} \lim_{m \uparrow \infty} \bar{c}(m)/\underline{c}(m) &= \lim_{m \uparrow \infty} (m - 1 + h)/m \\ &= 1, \end{aligned}$$

so as $m \uparrow \infty$, $c(m)/\underline{c}(m) \rightarrow 1$, and the continuous differentiability and strict concavity of $c(m)$ therefore implies

$$\lim_{m \uparrow \infty} c'(m) = \underline{c}'(m) = \bar{c}'(m) = \underline{\kappa}$$



Figure 5 Limiting MPC's

because any other fixed limit would eventually lead to a level of consumption either exceeding $\bar{c}(m)$ or lower than $\underline{c}(m)$.

Figure 5 confirms these limits visually. The top plot shows the converged consumption function along with its upper and lower bounds, while the lower plot shows the marginal propensity to consume.

Next we establish the limit of the expected consumption growth factor as $m_t \uparrow \infty$:

$$\lim_{m_t \uparrow \infty} \mathbb{E}_t[\mathbf{c}_{t+1}/\mathbf{c}_t] = \lim_{m_t \uparrow \infty} \mathbb{E}_t[\Gamma_{t+1}c_{t+1}/c_t].$$

But

$$\mathbb{E}_t[\Gamma_{t+1}\underline{c}_{t+1}/\bar{c}_t] \leq \mathbb{E}_t[\Gamma_{t+1}c_{t+1}/c_t] \leq \mathbb{E}_t[\Gamma_{t+1}\bar{c}_{t+1}/\underline{c}_t]$$

and

$$\lim_{m_t \uparrow \infty} \Gamma_{t+1}\underline{c}(m_{t+1})/\bar{c}(m_t) = \lim_{m_t \uparrow \infty} \Gamma_{t+1}\bar{c}(m_{t+1})/\underline{c}(m_t) = \lim_{m_t \uparrow \infty} \Gamma_{t+1}m_{t+1}/m_t,$$

while (for convenience defining $a(m_t) = m_t - c(m_t)$),

$$\begin{aligned} \lim_{m_t \uparrow \infty} \Gamma_{t+1}m_{t+1}/m_t &= \lim_{m_t \uparrow \infty} \left(\frac{\mathbf{R}a(m_t) + \Gamma_{t+1}\xi_{t+1}}{m_t} \right) \\ &= (\mathbf{R}\beta)^{1/\rho} = \mathbf{p} \end{aligned} \tag{48}$$



Figure 6 Upper and Lower Bounds on The Consumption Function

because $\lim_{m_t \uparrow \infty} a'(m) = \mathbf{P}_R$ ³⁹ and $\Gamma_{t+1}\xi_{t+1}/m_t \leq (\Gamma\bar{\psi}\bar{\theta}/(1 - \wp))/m_t$ which goes to zero as m_t goes to infinity.

Hence we have

$$\mathbf{P} \leq \lim_{m_t \uparrow \infty} \mathbb{E}_t[\mathbf{c}_{t+1}/\mathbf{c}_t] \leq \mathbf{P}$$

so as cash goes to infinity, consumption growth approaches its value \mathbf{P} in the perfect foresight model.

3.2 Limits as m Approaches Zero

Equation (41) shows that the limiting value of $\bar{\kappa}$ is

$$\bar{\kappa} = 1 - \mathbf{R}^{-1}(\wp \mathbf{R} \beta)^{1/\rho}.$$

Defining $e(m) = c(m)/m$ as before we have

$$\lim_{m \downarrow 0} e(m) = (1 - \wp^{1/\rho} \mathbf{P}_R) = \bar{\kappa}.$$

Now using the continuous differentiability of the consumption function along with

³⁹This is because $\lim_{m_t \uparrow \infty} a(m_t)/m_t = 1 - \lim_{m_t \uparrow \infty} c(m_t)/m_t = 1 - \lim_{m_t \uparrow \infty} c'(m_t) = \mathbf{P}_R$.

L'Hôpital's rule, we have

$$\lim_{m \downarrow 0} c'(m) = \lim_{m \downarrow 0} e(m) = \bar{\kappa}.$$

Figure 5 confirms that the numerical solution obtains this limit for the MPC as m approaches zero.

For consumption growth, as $m \downarrow 0$ we have

$$\begin{aligned} \lim_{m_t \downarrow 0} \mathbb{E}_t \left[\left(\frac{c(m_{t+1})}{c(m_t)} \right) \Gamma_{t+1} \right] &> \lim_{m_t \downarrow 0} \mathbb{E}_t \left[\left(\frac{\underline{c}(\mathcal{R}_{t+1}a(m_t) + \xi_{t+1})}{\bar{\kappa}m_t} \right) \Gamma_{t+1} \right] \\ &= \wp \lim_{m_t \downarrow 0} \mathbb{E}_t \left[\left(\frac{\underline{c}(\mathcal{R}_{t+1}a(m_t))}{\bar{\kappa}m_t} \right) \Gamma_{t+1} \right] \\ &\quad + (1 - \wp) \lim_{m_t \downarrow 0} \mathbb{E}_t \left[\left(\frac{\underline{c}(\mathcal{R}_{t+1}a(m_t) + \theta_{t+1}/(1 - \wp))}{\bar{\kappa}m_t} \right) \Gamma_{t+1} \right] \\ &> (1 - \wp) \lim_{m_t \downarrow 0} \mathbb{E}_t \left[\left(\frac{\underline{c}(\theta_{t+1}/(1 - \wp))}{\bar{\kappa}m_t} \right) \Gamma_{t+1} \right] \\ &= \infty \end{aligned}$$

where the second-to-last line follows because $\lim_{m_t \downarrow 0} \mathbb{E}_t \left[\left(\frac{\underline{c}(\mathcal{R}_{t+1}a(m_t))}{\bar{\kappa}m_t} \right) \Gamma_{t+1} \right]$ is positive, and the last line follows because the minimum possible realization of θ_{t+1} is $\underline{\theta} > 0$ so the minimum possible value of expected next-period consumption is positive.⁴⁰

3.3 Unique ‘Stable’ Points

Two theorems, whose substance is described here and whose details are in an appendix, articulate alternative (but closely related) stability criteria for the model.

3.3.1 ‘Individual Target Wealth’

One definition of a ‘stable’ point is an \hat{m} such that if $m_t = \hat{m}$, then $\mathbb{E}_t[m_{t+1}] = m_t$. Existence of such a target turns out to require the GIC-Nrm condition.

Theorem 2. *For the nondegenerate solution to the problem defined in Section 2.1 when FVAC, WRIC, and GIC-Nrm all hold, there exists a unique cash-on-hand-to-permanent-income ratio $\hat{m} > 0$ such that*

$$\mathbb{E}_t[m_{t+1}/m_t] = 1 \text{ if } m_t = \hat{m}. \quad (49)$$

Moreover, \hat{m} is a point of ‘wealth stability’ in the sense that

$$\begin{aligned} \forall m_t \in (0, \hat{m}), \quad \mathbb{E}_t[m_{t+1}] &> m_t \\ \forall m_t \in (\hat{m}, \infty), \quad \mathbb{E}_t[m_{t+1}] &< m_t. \end{aligned} \quad (50)$$

⁴⁰None of the arguments in either of the two prior sections depended on the assumption that the consumption functions had converged. With more cumbersome notation, each derivation could have been replaced by the corresponding finite-horizon versions. This strongly suggests that it should be possible to extend the circumstances under which the problem can be shown to define a contraction mapping to the union of the parameter values under which {RIC, FHCW} hold and {FVAC, WRIC} hold. That extension is not necessary for our purposes here, so we leave it for future work.

Since $m_{t+1} = (m_t - c(m_t))\mathcal{R}_{t+1} + \xi_{t+1}$, the implicit equation for \hat{m} is

$$\begin{aligned}\mathbb{E}_t[(\hat{m} - c(\hat{m}))\mathcal{R}_{t+1} + \xi_{t+1}] &= \hat{m} \\ (\hat{m} - c(\hat{m}))\underbrace{\mathcal{R} \mathbb{E}_t[\psi^{-1}]}_{\equiv \bar{\mathcal{R}}} + 1 &= \hat{m}\end{aligned}\tag{51}$$

3.3.2 Collective Stability and the Expected-Balanced-Growth State

A traditional question in macroeconomic models is whether there is a ‘balanced growth’ equilibrium in which aggregate variables (income, consumption, market resources) all grow forever at the same rate. For the model here, we have already seen in Figure 4 that there is no single m for which $\mathbb{E}_t[\mathbf{m}_{t+1}]/\mathbf{m}_t = \mathbb{E}_t[\mathbf{c}_{t+1}]/\mathbf{c}_t = \Gamma$ for an individual consumer. Nevertheless, analysis below will show that economies populated by collections of such consumers can exhibit balanced growth in the aggregate.

As an input to that analysis, we show here that if the GIC holds, the problem will have a point of what we call ‘collective stability’ by which we mean that there is some \check{m} such that, for all $m_t > \check{m}$, $\mathbb{E}_t[m_{t+1}/m_t] < \Gamma$, and conversely for $m_t < \check{m}$. (‘Collective’ is meant to capture the fact that calculating the expectation of the levels of future \mathbf{m} and \mathbf{p} before dividing \mathbf{m} by \mathbf{p} is akin to examining aggregate values in a population).

The critical m will be the value at which \mathbf{m} growth matches Γ :

$$\begin{aligned}\mathbb{E}_t[\mathbf{m}_{t+1}]/\mathbf{m}_t &= \mathbb{E}_t[\mathbf{p}_{t+1}]/\mathbf{p}_t \\ \mathbb{E}_t[m_{t+1}\Gamma\psi_{t+1}\mathbf{p}_t]/(m_t\mathbf{p}_t) &= \mathbb{E}_t[\mathbf{p}_t\Gamma\psi_{t+1}]/\mathbf{p}_t \\ \mathbb{E}_t\left[\psi_{t+1}\underbrace{((m_t - c(m_t))\mathcal{R}/(\Gamma\psi_{t+1})) + \xi_{t+1}}_{m_{t+1}}\right]/m_t &= 1 \\ \mathbb{E}_t\left[(\check{m} - c(\check{m}))\underbrace{\mathcal{R}}_{\bar{\mathcal{R}}/\Gamma} + \psi_{t+1}\xi_{t+1}\right] &= \check{m} \\ (\check{m} - c(\check{m}))\mathcal{R} + 1 &= \check{m}\end{aligned}\tag{52}$$

The only difference between (52) and (51) is the substitution of \mathcal{R} for $\bar{\mathcal{R}}$.

We will refer to \check{m} as the problem’s Expected-Balanced-Growth State, a term motivated by the fact that an economy composed of consumers all of whom had $m_t = \check{m}$, *would* exhibit balanced growth if all consumers happened to continually draw permanent and transitory shocks equal to their expected values of 1.0 forever.

Theorem 3 formally states the relevant proposition.

Theorem 3. *For the nondegenerate solution to the problem defined in Section 2.1 when FVAC, WRIC, and GIC all hold, there exists a unique pseudo-steady-state cash-on-hand-to-income ratio $\check{m} > 0$ such that*

$$\mathbb{E}_t[\psi_{t+1}m_{t+1}/m_t] = 1 \text{ if } m_t = \check{m}.\tag{53}$$

Moreover, \check{m} is a point of stability in the sense that

$$\forall m_t \in (0, \check{m}), \mathbb{E}_t[\mathbf{m}_{t+1}]/\mathbf{m}_t > \Gamma\tag{54}$$

$$\forall m_t \in (\check{m}, \infty), \mathbb{E}_t[\mathbf{m}_{t+1}]/\mathbf{m}_t < \Gamma.$$

The proofs of the two theorems are almost completely parallel; to save space, they are relegated to Appendix M. In sum, they involve three steps:

1. Existence and continuity of $\mathbb{E}_t[m_{t+1}/m_t]$ or $\mathbb{E}_t[m_{t+1}\psi_{t+1}/m_t]$
 - This follows from existence and continuity of the constituents
2. Existence of the equilibrium point
 - This follows from the upper and lower bound limiting MPC's, existence and continuity, and the Intermediate Value Theorem
3. Monotonicity of $\mathbb{E}_t[m_{t+1} - m_t]$ or $\mathbb{E}_t[m_{t+1}\psi_{t+1} - m_t]$
 - This follows from concavity of the consumption function

3.3.3 Example Where There Is An Expected-Balanced-Growth \check{m} State But No Target

Because the equations defining target and pseudo-steady-state m , (51) and (52), differ only by substitution of \mathcal{R} for $\bar{\mathcal{R}} = \mathcal{R}\mathbb{E}[\psi^{-1}]$, if there are no permanent shocks ($\psi \equiv 1$), the conditions are identical. For many parameterizations (e.g., under the baseline parameter values used for constructing figure 4), \hat{m} and \check{m} will not differ much.

An illuminating exception is exhibited in Figure 7, which modifies the baseline parameter values by quadrupling the variance of the permanent shocks, enough to cause failure of the GIC-Nrm; now there is no target wealth level \hat{m} ($c(m)$ remains everywhere below the level that would keep expected m constant).

The pseudo-steady-state still exists because it turns off realizations of the permanent shock. But an aggregate balanced growth equilibrium can exist even when realizations of the permanent shock are not turned off, and instead implemented exactly as specified in the model. The key insight can be understood by considering the evolution of an economy that starts, at date t , with the entire population at $m_t = \check{m}$, but then evolves according to the model's assumed dynamics between t and $t+1$. Equation (52) will still hold, so for this first period, at least, the economy will exhibit balanced growth: the growth factor for aggregate \mathbf{m} will match the growth factor for permanent income Γ . It is true that there will be people for whom $b_{t+1} = a_t \mathbf{R}/(\Gamma\psi_{t+1})$ is boosted by a small draw of ψ_{t+1} . But their contribution to the aggregate variable is given by $\mathbf{b}_{t+1} = b_{t+1}\psi_{t+1}$, so their \mathbf{b}_{t+1} is reweighted by an amount that exactly unwinds that boosting.

The surprising consequence is that, if the GIC holds but the GIC-Nrm fails, it is possible to construct an aggregate economy composed of consumers all of whom have target wealth of $\hat{m} = \infty$, but in which the aggregate economy still exhibits balanced growth with a finite ratio of aggregate wealth to income.

This is a good introduction to a more explicit discussion of aggregation.




Figure 7 {FVAC,GIC,~~GIC-Nrm~~}: No \hat{m} Exists But \check{m} Does

4 The Aggregate and Idiosyncratic Relationship Between Consumption Growth and Income Growth

Until now for convenience we have assumed infinite horizons, with the implicit understanding that Poisson mortality could be handled by adjusting the effective discount factor for mortality. On that basis, Section 4.1 continues to omit mortality. But a reason for explicitly introducing mortality will appear at the end of Section 4.2, so implications of alternative assumptions about mortality are briefly examined in Section 4.3.

Formally, assume a continuum of *ex ante* identical households on the unit interval, with constant total mass normalized to one and indexed by $i \in [0, 1]$, all behaving according to the model specified above. Szeidl (?) proves that whenever the GIC holds such a population will be characterized by invariant distributions of m , c , and a ;⁴¹ designate these \mathcal{F}^m , \mathcal{F}^a , and \mathcal{F}^c .

⁴¹?’s equation (9), in our notation, is:

$$\begin{aligned}\mathbb{E} \log \mathbf{R}(1 - \kappa) &< \mathbb{E} \log \Gamma \psi \\ \mathbb{E} \log \mathbf{R} \mathbf{P}_{\mathbf{R}} &< \mathbb{E} \log \Gamma \psi \\ \log \mathbf{P}_{\Gamma} &< \mathbb{E} \log \psi\end{aligned}$$

and under our assumption that $\log \psi \sim \mathcal{N}(-\sigma_{\psi}^2/2, \sigma_{\psi}^2)$ we can exponentiate both sides to obtain the GIC, $\mathbf{P}_{\Gamma} < 1$. If the permanent income shocks are not lognormally distributed, the expression must be tested in Szeidl’s original form.

4.1 Consumption and Income Growth at the Household Level

The operator $\mathbb{M}[\bullet]$ yields the mean of its argument in the population, as distinct from the expectations operator $\mathbb{E}[\bullet]$ used above, which represents beliefs about the future.

In a converged economy, an economist with a microeconomic dataset could calculate the average consumption growth rate as:

$$\begin{aligned}\mathbb{M}[\Delta \log \mathbf{c}_{t+1}] &= \mathbb{M}[\log c_{t+1} \mathbf{p}_{t+1} - \log c_t \mathbf{p}_t] \\ &= \mathbb{M}[\log \mathbf{p}_{t+1} - \log \mathbf{p}_t] + \mathbb{M}[\log c_{t+1} - \log c_t] \\ &= (\gamma - \sigma_\psi^2/2) + \mathbb{M}[\log c_{t+1} - \log c_t] \\ &= (\gamma - \sigma_\psi^2/2)\end{aligned}$$

where $\gamma = \log \Gamma$ and the last equality follows because the invariance of \mathcal{F}^c (?) means that $\mathbb{M}[\log c_{t+n}] = \mathbb{M}[\log c_t]$. Thus, the same GIC that guaranteed the existence of an ‘individual pseudo-steady-state’ value of m at the microeconomic level guarantees both that there will be an invariant distribution of the population across values of the model variables and that in that invariant distribution the mean growth rates of all idiosyncratic variables are the same (see ? for details).

4.2 Balanced Growth of Aggregate Income, Consumption, and Wealth

Using boldface capital letters for aggregates, the growth factor for aggregate income is:

$$\begin{aligned}\mathbf{Y}_{t+1}/\mathbf{Y}_t &= \mathbb{M}[\xi_{t+1} \Gamma \psi_{t+1} \mathbf{p}_t] / \mathbb{M}[\mathbf{p}_t \xi_t] \\ &= \Gamma\end{aligned}$$

because of the independence assumptions we have made about ξ and ψ .

From the perspective of period t ,

$$\begin{aligned}\mathbf{A}_{t+1} &= \mathbb{M}[a_{t+1} \mathbf{p}_{t+1}] \\ &= \Gamma \mathbb{M}[(a_t + (a_{t+1} - a_t)) \mathbf{p}_t \psi_{t+1}] \\ &= \Gamma \left(\underbrace{\mathbb{M}[a_t \mathbf{p}_t \psi_{t+1}]}_{=\mathbf{A}_t} + \underbrace{\mathbb{M}[(a_{t+1} - a_t)] \mathbb{M}[\mathbf{p}_t \psi_{t+1}]}_{=0 \text{ (?)}} + \text{cov}_t(a_{t+1} - a_t, \mathbf{p}_t \psi_{t+1}) \right) \\ \mathbf{A}_{t+1}/\mathbf{A}_t &= \Gamma \left(1 + \frac{\text{cov}(a_{t+1}, \mathbf{p}_t \psi_{t+1})}{\mathbb{M}[a_t \mathbf{p}_t]} \right)\end{aligned}$$

Unfortunately, the covariance term in the numerator, while generally small, will not in general be zero. This is because the realization of the permanent shock ψ_{t+1} has a nonlinear effect on a_{t+1} . Matters are simpler if there are no permanent shocks; see Appendix F for a proof that in that case the growth rate of assets (and other variables) does eventually converge to the growth rate of aggregate permanent income.

One way of thinking about the problem here is that it may reflect the fact that, under our assumptions, permanent income \mathbf{p} does not have an ergodic distribution; its distribution of becomes forever wider over time, because our consumers never die and each immortal person is perpetually subject to symmetric shocks to their $\log \mathbf{p}$.

This is why we need to introduce mortality.

4.3 Mortality and Redistribution

Most heterogeneous-agent models incorporate a constant positive probability of death, following ?. In such a Blanchardian model, for probabilities of death that exceed a threshold that depends on the size of the permanent shocks, ? show that the limiting distribution of permanent income has a finite variance, which is a useful step in the direction of taming the problems caused by an unbounded distribution of \mathbf{p} . Numerical results in that paper confirm the intuition that, under appropriate impatience conditions, balanced growth arises (though a formal proof remains elusive).

Even with those (numerical) results in hand, the centrality of mortality assumptions to the existence and nature of steady states requires them to be discussed briefly here.

4.3.1 Blanchard Lives

? assumes the existence of an annuitization scheme in which estates of dying consumers are redistributed to survivors in proportion to survivors' wealth, giving the recipients a higher effective rate of return. This treatment has several analytical advantages, most notably that the effect of mortality on the time preference factor is the exact inverse of its effect on the (effective) interest factor: If the probability of remaining alive is \aleph , then assuming that no utility accrues after death makes the effective discount factor $\hat{\beta} = \beta\aleph$, while the enhancement to the rate of return from the annuity scheme yields an effective interest rate of \hat{R}/\aleph (recall that because of Poisson mortality, the average wealth of the two groups is identical). Combining these, the effective patience factor in the new economy $\hat{\mathbf{P}}$ is unchanged from its value in the infinite horizon model:

$$\hat{\mathbf{P}} \equiv (\beta\aleph R/\aleph)^{1/\rho} = (R\beta)^{1/\rho} \equiv \mathbf{P}. \quad (55)$$

The only adjustments this requires to the analysis from prior parts of this paper are therefore to the few elements that involve a role for R distinct from its contribution to \mathbf{P} (principally, the RIC).

The numerical finding that the covariance term above is approximately zero allows us to conclude again that the key requirement for aggregate balanced growth is presumably the GIC.

4.3.2 Modigliani Lives

?’s innovation was useful not only for the insight it provided but also because the principal alternative, the Life Cycle model of ?, was computationally challenging given the then-available technologies. Aside from its (considerable) conceptual value, there is no need for Blanchard’s analytical solution today, when serious modeling incorporates uncertainty, constraints, and other features that rule out analytical solutions anyway.

The simplest alternative to Blanchard’s mortality is to follow Modigliani in assuming that any wealth remaining at death occurs accidentally (not implausible, given the robust

finding that for the great majority of households, bequests amount to less than 2 percent of lifetime earnings, ??).

Even if bequests are accidental, a macroeconomic model must make some assumption about how they are disposed of: As windfalls to heirs, estate tax proceeds, etc. We again consider the simplest choice, because it again represents something of a polar alternative to Blanchard: Without a bequest motive, there are no behavioral effects of a 100 percent estate tax; we assume such a tax is imposed and that the revenues are effectively thrown in the ocean; the estate-related wealth effectively vanishes from the economy.

The chief appeal of this approach is the simplicity of the change it makes in the condition required for the economy to exhibit a balanced growth equilibrium. If \aleph is the probability of remaining alive, the condition changes from the plain GIC to a looser mortality-adjusted GIC:

$$\aleph \mathfrak{D}_T < 1. \tag{56}$$

With no income growth, the condition required to prohibit unbounded growth in aggregate wealth would be the condition that prevents the per-capita wealth to income ratio of surviving consumers from growing faster than the rate at which mortality diminishes their collective population. With income growth, the aggregate wealth-to-income ratio will head to infinity only if a cohort of consumers is patient enough to make the desired rate of growth of wealth fast enough to counteract combined erosive forces of mortality and productivity.

5 Conclusions

Numerical solutions to optimal consumption problems, in both life cycle and infinite horizon contexts, have become standard tools since the first reasonably realistic models were constructed in the late 1980s. One contribution of this paper is to show that finite horizon (‘life cycle’) versions of the simplest such models, with assumptions about income shocks (transitory and permanent) dating back to ? and standard specifications of preferences — and without plausible (but computationally and mathematically inconvenient) complications like liquidity constraints — have attractive properties (like continuous differentiability of the consumption function, and analytical limiting MPC’s as resources approach their minimum and maximum possible values).

The main focus of the paper, though, is on the limiting solution of the finite horizon model as the horizon extends to infinity. The paper shows that the simple model has additional attractive properties: A ‘**Finite Value of Autarky**’ condition guarantees convergence of the consumption function, under the mild additional requirement of a ‘Weak Return Impatience Condition’ that will never bind for plausible parameterizations, but provides intuition for the bridge between this model and models with explicit liquidity constraints. The paper also provides a roadmap for the model’s relationships to the perfect foresight model without and with constraints. The constrained perfect foresight model provides an upper bound to the consumption function (and value function) for the model with uncertainty, which explains why the conditions for the model to

have a nondegenerate solution closely parallel those required for the perfect foresight constrained model to have a nondegenerate solution.

The main use of infinite horizon versions of such models is in heterogeneous agent macroeconomics. The paper articulates intuitive ‘**Growth Impatience Conditions**’ under which populations of such agents, with Blanchardian (tighter) or Modiglianian (looser) mortality will exhibit balanced growth. Finally, the paper provides the analytical basis for many results about buffer-stock saving models that are so well understood that even without analytical foundations researchers uncontroversially use them as explanations of real-world phenomena like the cross-sectional pattern of consumption dynamics in the Great Recession.